

NLP and Sentiment Analysis

Name: Cassidy Exum

Date: 03/13/2023

Using data from CrowdFlower, I'm going to create different NLP models and perform sentiment analysis on tweets talking about Apple and Google.

Reminder - This is a classification task

Preliminary thoughts on process -

1. pre-preparation - vectorize all tweets at start
2. Classify all tweets as positive, negative, neutral
3. exploration - Begin exploring products that are discussed positive or negatively and check word associativity to determine the reasons for those sentiments.
4. Give advertising recommendations - Advertising recommendations: Marketing research shows we should advertise to people who like our products. If twitter is generally positive, suggest more money spent on advertising, if twitter is generally negative, perhaps less money should be spent. Or considering advertising that highlights the best aspects of the discussed products, and counters/discusses the improvements of negatively talked about products.
5. Give product development recommendations - Overwhelmingly negatively talked about products should be improved based on user feedback, determine most discussed feedback

Blog Post recapping this project and my thoughts:

<https://exumexaminesdata.blogspot.com/2023/03/nlp-and-sentiment-analysis-of-tech.html>

Summary

For this analysis I used F1 score as my evaluation metric, this score takes both false positives and false negatives into account. It is suitable for uneven class distribution problems.

I began by investigating the data. From the link provided for downloading the data we know it comes from around 2011 and 2013. Unfortunately the data only has 9000 tweets and the class imbalance is very prevalent with 60% of the tweets being "No emotion" which is a hinderance for training. A large amount of tweets were also considered to not be aimed at any particular brand or company, which I found to be untrue. I didn't want to mess up anything in the data and decided to leave it as is instead of imputing what tweets were aimed at different brands based on words used in the tweet.

For feature generation I tried a bunch of different things, including number of sentences, if a tweet contains an emoji, and vectorizers. The first batch of models I tried included multinomialNB and an SVC. The SVC had better accuracy and F1, but had a generally poor F1

score of only 58%. I then tried using logistic regression, random forest, and XGBoost. These models showed increased performance. I then went on to create a perceptron neural net using Keras. This model performed the best, with a 74% accuracy on the training set and 66% on the validation set. I created a new random validation set of 1800 rows and the neural net obtained an F1 score of 75%. Before declaring this the best model I decided to try a Word2Vec approach using the Random Forest and XGB because they performed well previously. The random forest and XGB had accuracy's of 67% and 68% respectively, but had F1 scores of 95% and 93%. The validation set was created again using 1800 randomly sampled rows from the training set.

The accuracy and F1 of the Neural Net is incredibly impressive and makes me think there was no overfitting. The F1 score is almost too good to be true when talking about the random forest and XGB models.

Resources

These notebooks and sites were a useful tools and references in completing this analysis

1. <https://www.kaggle.com/code/tanulsingh077/twitter-sentiment-extaction-analysis-eda-and-model>
2. <https://www.kaggle.com/code/nitin194/twitter-sentiment-analysis-word2vec-doc2vec>
3. <https://www.kaggle.com/code/prakharrathi25/sentiment-analysis-using-bert>
4. <https://github.com/learn-co-curriculum/dsc-classification-with-word-embeddings-codealong>

Preprocessing and Feature Engineering

```
In [1]: # imports

import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
from sklearn.model_selection import train_test_split
from matplotlib.ticker import MaxNLocator
import seaborn as sns

# NLTK
from nltk.tokenize import RegexpTokenizer
from nltk import FreqDist
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.naive_bayes import MultinomialNB
from sklearn.model_selection import cross_val_score
from nltk.corpus import stopwords

from nltk.stem.snowball import SnowballStemmer
from nltk.tokenize import sent_tokenize

#Sklearn
from sklearn.metrics import plot_confusion_matrix
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.pipeline import Pipeline
from sklearn.svm import SVC
```

```

from xgboost import XGBClassifier
from sklearn.metrics import f1_score
from sklearn.metrics import confusion_matrix

# warnings
import warnings
warnings.filterwarnings('ignore')

%matplotlib inline

```

```
In [2]: stopwords_list = stopwords.words('english')
```

```
In [3]: # create df with data and inspect
df = pd.read_csv('Data/data.csv', engine='python')
df.head()
```

```
Out[3]:
```

| | tweet_text | emotion_in_tweet_is_directed_at | is_there_an_emotion_directed_at_a_brand_or_pro |
|---|---|---------------------------------|--|
| 0 | .@wesley83 I have a 3G iPhone. After 3 hrs twe... | iPhone | Negative emotion |
| 1 | @jessedee Know about @fludapp ? Awesome iPad/i... | iPad or iPhone App | Positive emotion |
| 2 | @swonderlin Can not wait for #iPad 2 also. The... | iPad | Positive emotion |
| 3 | @sxsw I hope this year's festival isn't as cra... | iPad or iPhone App | Negative emotion |
| 4 | @sxtxstate great stuff on Fri #SXSW: Marissa M... | Google | Positive emotion |

```
In [4]: # Change column names for ease of use
df.columns = ['tweet', 'subject', 'emotion']
```

```
In [5]: # Check emotion distribution
df['emotion'].value_counts()
```

```
Out[5]:
```

| | |
|------------------------------------|------|
| No emotion toward brand or product | 5389 |
| Positive emotion | 2978 |
| Negative emotion | 570 |
| I can't tell | 156 |
| Name: emotion, dtype: int64 | |

```
In [6]: # Store any unknown emotion values elsewhere and then drop them from the main DF
# We can potentially use this later as a production version of the learner to gi
```

```
unknowns_df = df.loc[df['emotion'] == "I can't tell"]
```

```
In [7]: # Check NaNs
df.isna().sum()
```

```
Out[7]: tweet      1
subject    5655
emotion     0
dtype: int64
```

The number NaN's in the subject column is a bit alarming. I can't really do any simple Imputing to reduce that number, and I can drop over half my dataset. Either have to live with it and move on, or iterate through the df and see if the creators of the data missed things. Could perhaps just use an If/Then to check is if a group of apple or apple related words is in the tweet and assign that tweet to apple, same with google.

```
In [8]: #Drop NaN values only if tweet is NaN
df = df.dropna(subset=['tweet'])
df.isna().sum()
```

```
Out[8]: tweet      0
subject    5654
emotion     0
dtype: int64
```

```
In [9]: df.head()
```

```
Out[9]:
```

| | tweet | subject | emotion |
|---|---|--------------------|------------------|
| 0 | .@wesley83 I have a 3G iPhone. After 3 hrs twe... | iPhone | Negative emotion |
| 1 | @jessedee Know about @fludapp ? Awesome iPad/i... | iPad or iPhone App | Positive emotion |
| 2 | @swonderlin Can not wait for #iPad 2 also. The... | iPad | Positive emotion |
| 3 | @sxsw I hope this year's festival isn't as cra... | iPad or iPhone App | Negative emotion |
| 4 | @sxtxstate great stuff on Fri #SXSW: Marissa M... | Google | Positive emotion |

```
In [10]: # Lower case all tweets
df['tweet'] = df['tweet'].str.lower()
```

```
In [11]: df.subject.value_counts()
```

```
Out[11]: iPad      942
Apple      659
iPad or iPhone App  470
Google      429
iPhone      296
Other Google product or service  292
Android App    81
Android        78
Other Apple product or service  35
Name: subject, dtype: int64
```

```
In [12]: #investigate other google product tweets
google_rows = df[df['subject'] == 'Other Google product or service']
google_rows
```

Out[12]:

| | tweet | subject | emotion |
|------|---|---------------------------------|------------------|
| 13 | gotta love this #sxsw google calendar featurin... | Other Google product or service | Positive emotion |
| 27 | someone started an #austin @partnerhub group i... | Other Google product or service | Positive emotion |
| 163 | just left #sxsw tradeshow demo of @mention at ... | Other Google product or service | Positive emotion |
| 198 | sweet new 3-d google maps demo going on in bal... | Other Google product or service | Positive emotion |
| 199 | more than 35 million miles per day are driving... | Other Google product or service | Positive emotion |
| ... | ... | ... | ... |
| 8989 | it's crazy how much culture is documented in g... | Other Google product or service | Positive emotion |
| 8992 | looks very interesting rt@mention google to la... | Other Google product or service | Positive emotion |
| 9006 | creativity prompt: use google maps to virtuall... | Other Google product or service | Positive emotion |
| 9025 | absolutely! rt @mention timely good schtuff f... | Other Google product or service | Positive emotion |
| 9080 | diller says google tv "might be run over ... | Other Google product or service | Negative emotion |

292 rows x 3 columns

I think for now, its ok to just impute the missing values based on Apple, iPhone, iPad, Google, or Andriod, and not worry too much about the related products until later. Later I can use word association to find tweets with 'app' in them as well to further expand subject

In [13]:

```
# # Impute NaN values in subject because things were missed

# apple words
df.loc[df['tweet'].str.contains('apple', case=False) & df['subject'].isna(), 'subject'] = 'apple'
df.loc[df['tweet'].str.contains('iphone', case=False) & df['subject'].isna(), 'subject'] = 'iphone'
df.loc[df['tweet'].str.contains('ipad', case=False) & df['subject'].isna(), 'subject'] = 'ipad'
df.loc[df['tweet'].str.contains('itunes', case=False) & df['subject'].isna(), 'subject'] = 'itunes'

# google
df.loc[df['tweet'].str.contains('google', case=False) & df['subject'].isna(), 'subject'] = 'google'

# andriod
df.loc[df['tweet'].str.contains('andriod', case=False) & df['subject'].isna(), 'subject'] = 'andriod'
```

In [14]:

```
#check NaN again
df.isna().sum()
```

Out[14]:

```
tweet      0
subject    907
emotion     0
dtype: int64
```

That was a massive massive success. Lets move on to tokenizing and removing stopwords

```
In [15]: # Tokenize Tweets

basic_token_pattern = r"(?u)\b\w\w+\b"

tokenizer = RegexpTokenizer(basic_token_pattern)

df["tweet_tokenized"] = df["tweet"].apply(tokenizer.tokenize)

# Display full text
#df.style.set_properties(**{'text-align': 'left'})
```

```
In [16]: #remove stopwords

# this helper funtion returns a list with any stopwords in the original list removed
def remove_stopwords(token_list):
    """
    Given a list of tokens, return a list where the tokens
    that are also present in stopwords_list have been
    removed
    """
    return [w for w in token_list if w not in stopwords_list]

df["tweet_without_stopwords"] = df["tweet_tokenized"].apply(remove_stopwords)
```

I think now would be a good time to create frequency distributions for any Apple related tweets, google related tweets, and then again with positive, negative, and no emotion tweets


```
In [17]: def visualize_top_10(freq_dist, title):

    # Extract data for plotting
    top_10 = list(zip(*freq_dist.most_common(10)))
    tokens = top_10[0]
    counts = top_10[1]

    # Set up plot and plot data
    fig, ax = plt.subplots()
    ax.bar(tokens, counts)

    # Customize plot appearance
    ax.set_title(title)
    ax.set_ylabel("Count")
    ax.yaxis.set_major_locator(MaxNLocator(integer=True))
    ax.tick_params(axis="x", rotation=90)
```

```
In [127... entire_df_freq_dist = FreqDist(df["tweet_without_stopwords"].explode())
visualize_top_10(entire_df_freq_dist, "Top 10 Word Frequency")
```




It might be worth removing sxsw, mention, link, and rt as well. They don't give me any information.

1. #sxsw is the tech event most of these tweets are talking about
2. mention refers quoting someone
3. link refers to links in tweets
4. rt refers to a retweet
5. quot refers to a quote retweet
6. amp refers to an ampersand

```
In [19]: stopwords_list.append('sxsw')
stopwords_list.append('mention')
stopwords_list.append('link')
stopwords_list.append('rt')
stopwords_list.append('quot')
stopwords_list.append('amp')
df["tweet_without_stopwords"] = df["tweet_tokenized"].apply(remove_stopwords)
```

```
In [126... entire_df__freq_dist = FreqDist(df["tweet_without_stopwords"].explode())
visualize_top_10(entire_df__freq_dist, "Top 10 Word Frequency")
```



```
In [21]: df.subject.value_counts()
```

```
Out[21]: Google                2093
         iPad                 1916
         Apple                1841
         iPhone              1166
         iPad or iPhone App    527
         Other Google product or service 292
         Android App           81
         Android              78
         Other Apple product or service 35
         Name: subject, dtype: int64
```

```
In [22]: # Function that iterates through the different Subjects and plots their distribu

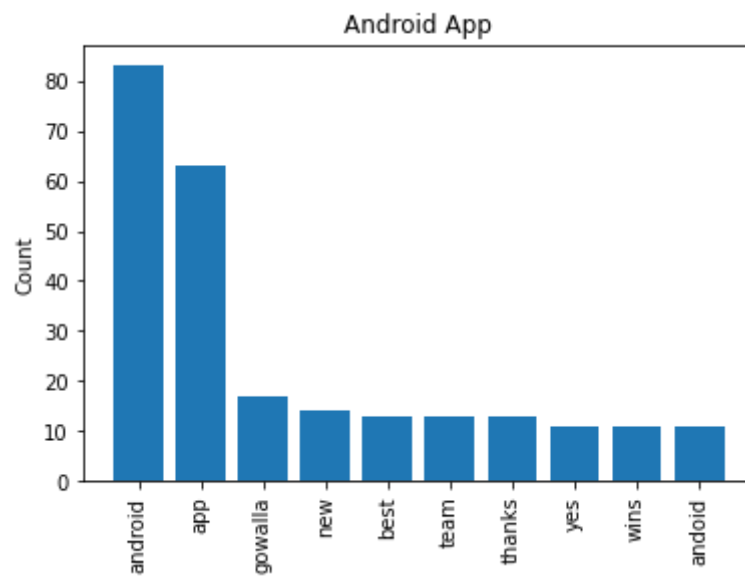

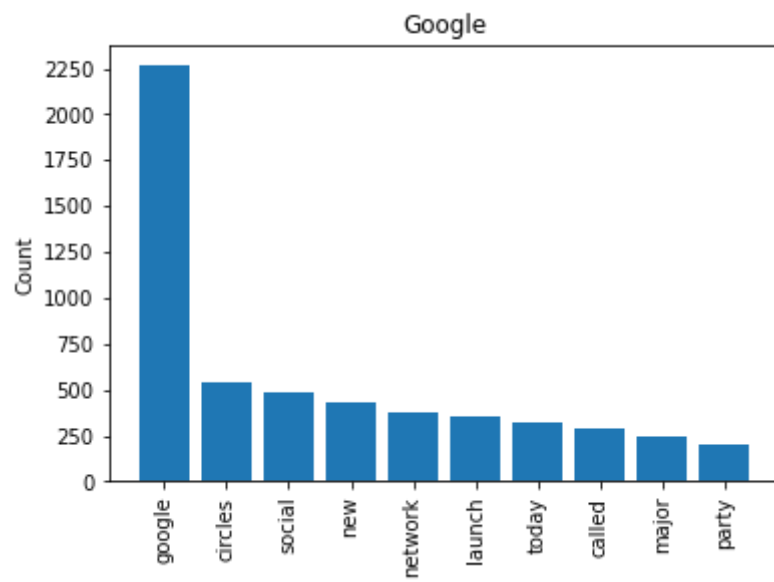
def freq_distribution_by_subject(df):
    #subjects = ['Google', 'iPad', 'Apple', 'iPhone', 'iPad or iPhone App']
    subjects = set(df['subject'].values.tolist())
    subjects.pop()


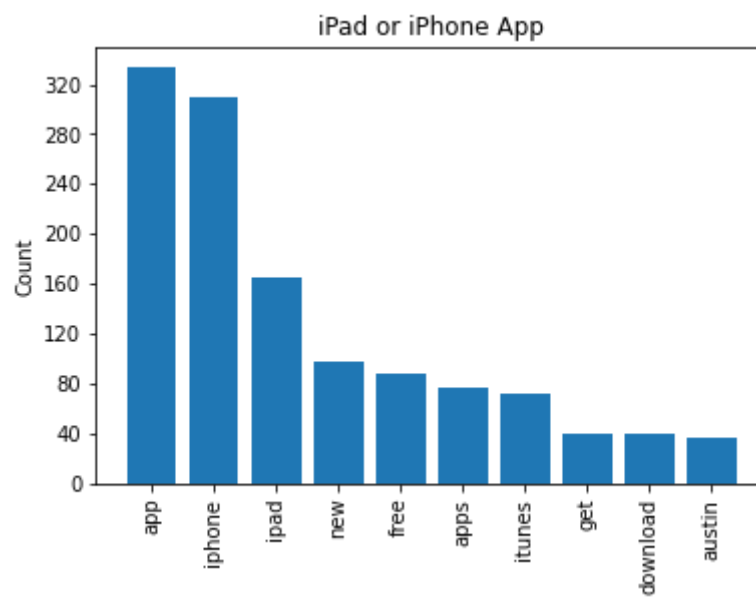
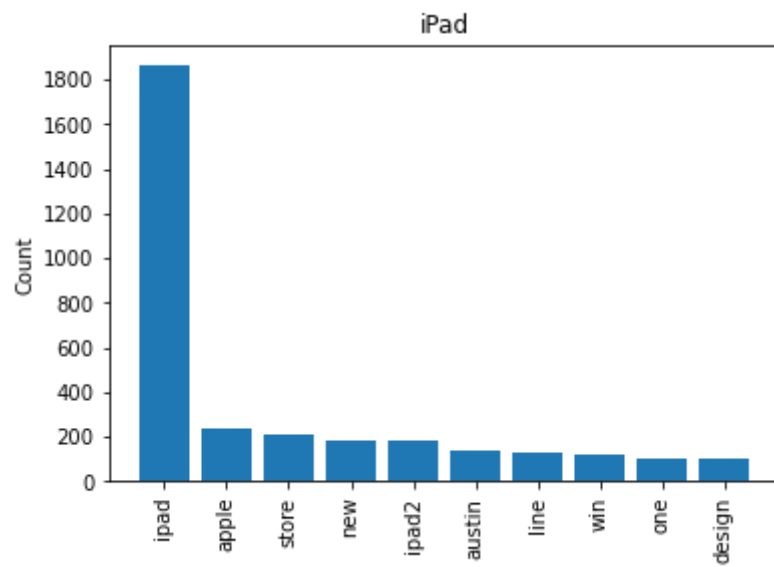
    #in the loop, create a freq dist for only a subject, then call the visualize

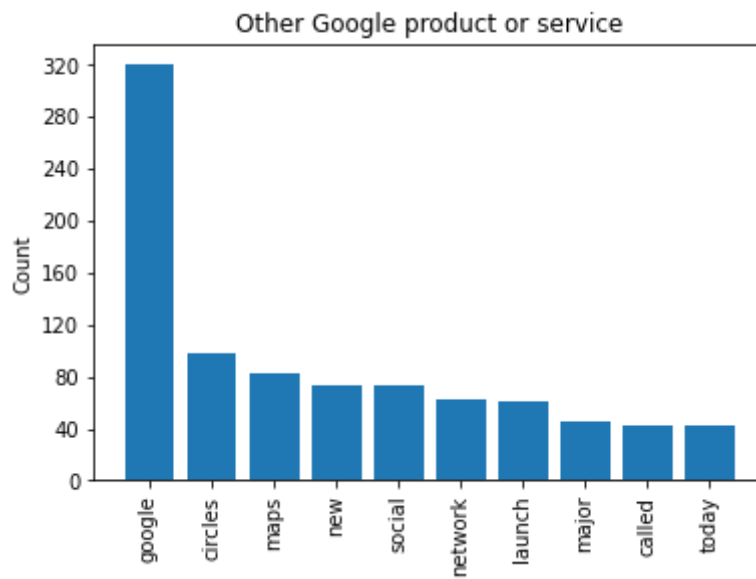
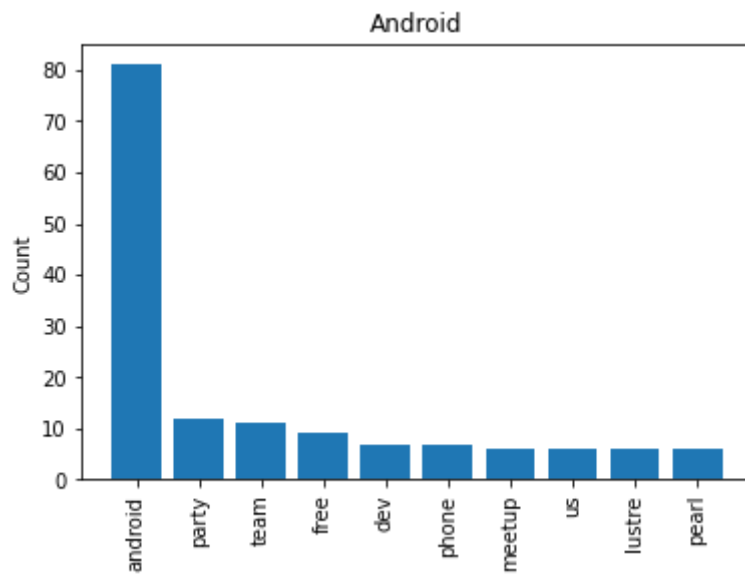
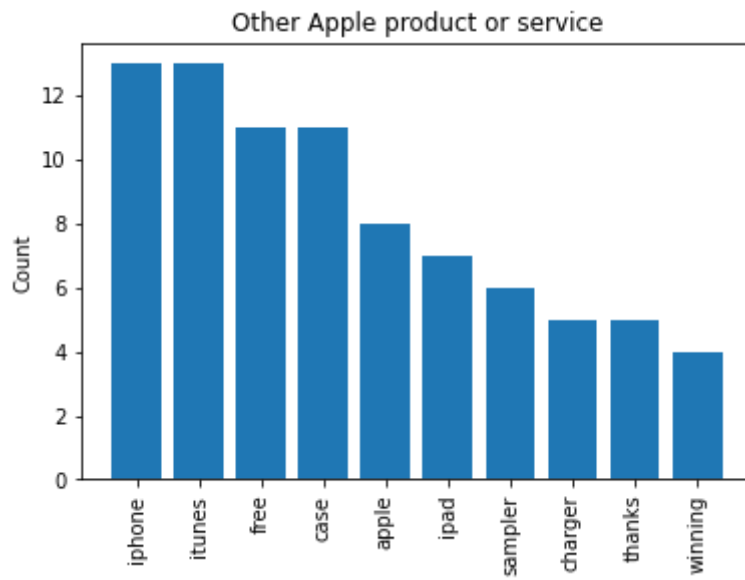
    for sub in subjects:
        print(sub)
        subject_df = df[df['subject'] == sub]
        subject_df_freq_dist = FreqDist(subject_df["tweet_without_stopwords"].ex
        visualize_top_10(subject_df_freq_dist, sub)
```

```
In [23]: freq_distribution_by_subject(df)
```

```
Google
Apple
Android App
iPad
iPad or iPhone App
iPhone
Other Apple product or service
Android
Other Google product or service
```





Most of these are probably not needed as the max word count is quite low, we will keep them for now.

Modeling

I'm going to start by using MultinomialNB and potentially iterating to a deep learning neural net to determine sentiment

```
In [24]: df.head()
```

| Out[24]: | tweet | subject | emotion | tweet_tokenized | tweet_without_stopwords |
|----------|---|--------------------|------------------|--|--|
| 0 | .@wesley83 i have a 3g iphone. after 3 hrs twe... | iPhone | Negative emotion | [wesley83, have, 3g, iphone, after, hrs, tweet...] | [wesley83, 3g, iphone, hrs, tweeting, rise_aus...] |
| 1 | @jessedee know about @fludapp ? awesome ipad/i... | iPad or iPhone App | Positive emotion | [jessedee, know, about, fludapp, awesome, ipad...] | [jessedee, know, fludapp, awesome, ipad, iphon...] |
| 2 | @swonderlin can not wait for #ipad 2 also. the... | iPad | Positive emotion | [swonderlin, can, not, wait, for, ipad, also, ...] | [swonderlin, wait, ipad, also, sale] |
| 3 | @sxsw i hope this year's festival isn't as cra... | iPad or iPhone App | Negative emotion | [sxsw, hope, this, year, festival, isn, as, cr...] | [hope, year, festival, crashy, year, iphone, app] |
| 4 | @sxtxstate great stuff on fri #sxsw: marissa m... | Google | Positive emotion | [sxtxstate, great, stuff, on, fri, sxsw, maris...] | [sxtxstate, great, stuff, fri, marissa, mayer,...] |

```
In [25]: # Start by setting up our training and test sets

y = df['emotion']
X = df['tweet']

X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=42, test_
X_train = pd.DataFrame(X_train)
X_test = pd.DataFrame(X_test)
y_train = pd.DataFrame(y_train)
y_test = pd.DataFrame(y_test)
```

```
In [26]: X_train.head()
```

| Out[26]: | tweet |
|----------|---|
| 3363 | having a great time at the google party #sxswi... |
| 3204 | @mention from:ubersocial for iphone now in the... |
| 4460 | are you in town for #sxsw? be sure to check in... |
| 2311 | the ironic tee has been usurped by the ipad 2 ... |
| 6298 | rt @mention marissa mayer: google will connect... |

```
In [27]: # Check distribution of training set
y_train.value_counts()
```

| | | |
|----------|------------------------------------|------|
| Out[27]: | emotion | |
| | No emotion toward brand or product | 4032 |

```
Positive emotion          2259
Negative emotion          411
dtype: int64
```

```
In [28]: # Check distribution of test set
y_test.value_counts()
```

```
Out[28]: emotion
No emotion toward brand or product    1356
Positive emotion                      719
Negative emotion                      159
dtype: int64
```

```
In [29]: #Create Baseline Model, create a TfidfVectorizer

tfidf = TfidfVectorizer(max_features=10)

# Fit the vectorizer on X_train["text"] and transform it
X_train_vectorized = tfidf.fit_transform(X_train['tweet'])

# Visually inspect the 10 most common words
pd.DataFrame.sparse.from_spmatrix(X_train_vectorized, columns=tfidf.get_feature_names_out())
```

```
Out[29]:
```

| | at | for | google | ipad | link | mention | rt | sxsw | tl |
|------|----------|----------|----------|----------|----------|----------|----------|----------|---------|
| 0 | 0.561447 | 0.000000 | 0.598490 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.260088 | 0.50886 |
| 1 | 0.000000 | 0.523624 | 0.000000 | 0.000000 | 0.388421 | 0.355803 | 0.000000 | 0.219420 | 0.42929 |
| 2 | 0.000000 | 0.823886 | 0.397220 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.172621 | 0.00000 |
| 3 | 0.320355 | 0.000000 | 0.000000 | 0.341491 | 0.000000 | 0.000000 | 0.000000 | 0.148403 | 0.87109 |
| 4 | 0.000000 | 0.000000 | 0.503358 | 0.000000 | 0.387227 | 0.354709 | 0.489748 | 0.218746 | 0.42797 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 6697 | 0.434081 | 0.000000 | 0.462720 | 0.000000 | 0.355965 | 0.326072 | 0.450209 | 0.402171 | 0.00000 |
| 6698 | 0.468943 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.704520 | 0.486367 | 0.217235 | 0.00000 |
| 6699 | 0.383120 | 0.000000 | 0.000000 | 0.000000 | 0.314175 | 0.287792 | 0.397356 | 0.177479 | 0.69447 |
| 6700 | 0.329455 | 0.000000 | 0.000000 | 0.351191 | 0.000000 | 0.742439 | 0.000000 | 0.152618 | 0.29859 |
| 6701 | 0.000000 | 0.761199 | 0.000000 | 0.000000 | 0.564652 | 0.000000 | 0.000000 | 0.318974 | 0.00000 |

6702 rows x 10 columns

```
In [30]: baseline_model = MultinomialNB()

# Evaluate the classifier on X_train_vectorized and y_train
baseline_cv = cross_val_score(baseline_model, X_train_vectorized, y_train['emotion'], cv=5)
baseline_cv
```

```
Out[30]: array([0.60178971, 0.601044 , 0.60149254, 0.60149254, 0.60223881])
```

```
In [31]: y_train.value_counts(normalize=True)
```

```
Out[31]: emotion
No emotion toward brand or product    0.601611
Positive emotion                      0.337064
```

```
Negative emotion                                0.061325
dtype: float64
```

Ok, so the baseline model is 60% accurate, but based on the distribution, if it only guess no emotion, thats what it would score, lets see if we can raise that value.

Start by Removing Stopwords

```
In [32]: #Removing Stopwords

tfidf = TfidfVectorizer(
    max_features=10,
    stop_words=stopwords_list
)

# Fit the vectorizer on X_train["text"] and transform it
X_train_vectorized = tfidf.fit_transform(X_train['tweet'])

# Visually inspect the vectorized data
pd.DataFrame.sparse.from_spmatrix(X_train_vectorized, columns=tfidf.get_feature_
```

```
Out[32]:
```

| | app | apple | austin | circles | google | ipad | iphone | new | social | store |
|------|----------|---------|----------|---------|--------|------|---------|-----|--------|----------|
| 0 | 0.000000 | 0.00000 | 0.000000 | 0.0 | 1.0 | 0.0 | 0.00000 | 0.0 | 0.0 | 0.000000 |
| 1 | 0.656654 | 0.00000 | 0.000000 | 0.0 | 0.0 | 0.0 | 0.52949 | 0.0 | 0.0 | 0.537072 |
| 2 | 0.000000 | 0.00000 | 0.000000 | 0.0 | 1.0 | 0.0 | 0.00000 | 0.0 | 0.0 | 0.000000 |
| 3 | 0.000000 | 0.00000 | 0.000000 | 0.0 | 0.0 | 1.0 | 0.00000 | 0.0 | 0.0 | 0.000000 |
| 4 | 0.000000 | 0.00000 | 0.000000 | 0.0 | 1.0 | 0.0 | 0.00000 | 0.0 | 0.0 | 0.000000 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 6697 | 0.000000 | 0.00000 | 0.000000 | 0.0 | 1.0 | 0.0 | 0.00000 | 0.0 | 0.0 | 0.000000 |
| 6698 | 0.000000 | 0.00000 | 0.000000 | 0.0 | 0.0 | 0.0 | 1.00000 | 0.0 | 0.0 | 0.000000 |
| 6699 | 0.000000 | 0.00000 | 0.000000 | 0.0 | 0.0 | 0.0 | 0.00000 | 0.0 | 0.0 | 0.000000 |
| 6700 | 0.000000 | 0.00000 | 0.000000 | 0.0 | 0.0 | 1.0 | 0.00000 | 0.0 | 0.0 | 0.000000 |
| 6701 | 0.000000 | 0.75127 | 0.500424 | 0.0 | 0.0 | 0.0 | 0.00000 | 0.0 | 0.0 | 0.430313 |

6702 rows x 10 columns

```
In [33]: stopwords_removed_cv = cross_val_score(baseline_model, X_train_vectorized, y_train, cv=5)
stopwords_removed_cv
```

```
Out[33]: array([0.60253542, 0.61073826, 0.60895522, 0.60671642, 0.59925373])
```

```
In [34]: print(baseline_cv.mean())
print(stopwords_removed_cv.mean())
```

```
0.6016115173572851
0.6056398098990506
```

Removing stopwords is very slightly better so we will keep them removed going forward. Lets now begin stemming words as well.

```
In [35]: stemmer = SnowballStemmer(language="english")

def stem_and_tokenize(document):
    tokens = tokenizer.tokenize(document)
    return [stemmer.stem(token) for token in tokens]

stemmed_stopwords = [stemmer.stem(word) for word in stopwords_list]
```

```
In [36]: tfidf = TfidfVectorizer(max_features=10,
                                stop_words=stemmed_stopwords,
                                tokenizer=stem_and_tokenize)

# Fit the vectorizer on X_train["text"] and transform it
X_train_vectorized = tfidf.fit_transform(X_train['tweet'])

# Visually inspect the vectorized data
pd.DataFrame.sparse.from_spmatrix(X_train_vectorized, columns=tfidf.get_feature_
```

```
Out[36]:
```

| | app | appl | austin | circl | googl | ipad | iphon | launch | new | store |
|------|---------|----------|----------|-------|-------|------|----------|--------|-----|----------|
| 0 | 0.00000 | 0.000000 | 0.000000 | 0.0 | 1.0 | 0.0 | 0.000000 | 0.0 | 0.0 | 0.000000 |
| 1 | 0.63101 | 0.000000 | 0.000000 | 0.0 | 0.0 | 0.0 | 0.545373 | 0.0 | 0.0 | 0.551720 |
| 2 | 0.00000 | 0.000000 | 0.000000 | 0.0 | 1.0 | 0.0 | 0.000000 | 0.0 | 0.0 | 0.000000 |
| 3 | 0.00000 | 0.000000 | 0.000000 | 0.0 | 0.0 | 1.0 | 0.000000 | 0.0 | 0.0 | 0.000000 |
| 4 | 0.00000 | 0.000000 | 0.000000 | 0.0 | 1.0 | 0.0 | 0.000000 | 0.0 | 0.0 | 0.000000 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 6697 | 0.00000 | 0.000000 | 0.000000 | 0.0 | 1.0 | 0.0 | 0.000000 | 0.0 | 0.0 | 0.000000 |
| 6698 | 0.00000 | 0.000000 | 0.000000 | 0.0 | 0.0 | 0.0 | 1.000000 | 0.0 | 0.0 | 0.000000 |
| 6699 | 0.00000 | 0.000000 | 0.000000 | 0.0 | 0.0 | 0.0 | 0.000000 | 0.0 | 0.0 | 0.000000 |
| 6700 | 0.00000 | 0.000000 | 0.000000 | 0.0 | 0.0 | 1.0 | 0.000000 | 0.0 | 0.0 | 0.000000 |
| 6701 | 0.00000 | 0.751687 | 0.500963 | 0.0 | 0.0 | 0.0 | 0.000000 | 0.0 | 0.0 | 0.428955 |

6702 rows × 10 columns

```
In [37]: stemmed_cv = cross_val_score(baseline_model, X_train_vectorized, y_train['emotio
stemmed_cv
```

```
Out[37]: array([0.59731544, 0.61222968, 0.60746269, 0.60223881, 0.59552239])
```

```
In [38]: print(baseline_cv.mean())
print(stopwords_removed_cv.mean())
print(stemmed_cv.mean())
```



```
0.6016115173572851
0.6056398098990506
0.6029537992364797
```


Well I didn't really expect that to do better after seeing it removed the letter e from many important words. Is sentence tokenizing valuable? Most tweets are single sentences regardless when capped at 140 or 280 characters...

```
In [39]: df["num_sentences"] = df["tweet"].apply(lambda x: len(sent_tokenize(x)))
```

```
In [40]: def freq_sent_distribution_by_emotion(df):  
    emotions = set(df['emotion'].values.tolist())  
  
    #in the loop, create a freq dist for only a subject, then call the visualize  
  
    for em in emotions:  
        emotion_df = df[df['emotion'] == em]  
        emotion_df_freq_dist = FreqDist(emotion_df["num_sentences"].explode())  
        visualize_top_10(emotion_df_freq_dist, em)
```

```
In [41]: freq_sent_distribution_by_emotion(df)
```





Ok maybe I'm wrong about what I expected, lets try creating another model and seeing if this can help.

```
In [42]: X_train["num_sentences"] = X_train['tweet'].apply(lambda x: len(sent_tokenize(x)))
X_train.num_sentences.value_counts()
```

```
Out[42]: 1    2911
         2    2312
         3    1133
         4     277
         5      62
         6       4
         7        3
         Name: num_sentences, dtype: int64
```

Lets try adding an emoticons deature as well

```
In [43]: # Emoticons

emoticon_query = r'(?::[\.;X=B][-^]?[)\]3D([OP/\||])(?:(?=\s))'


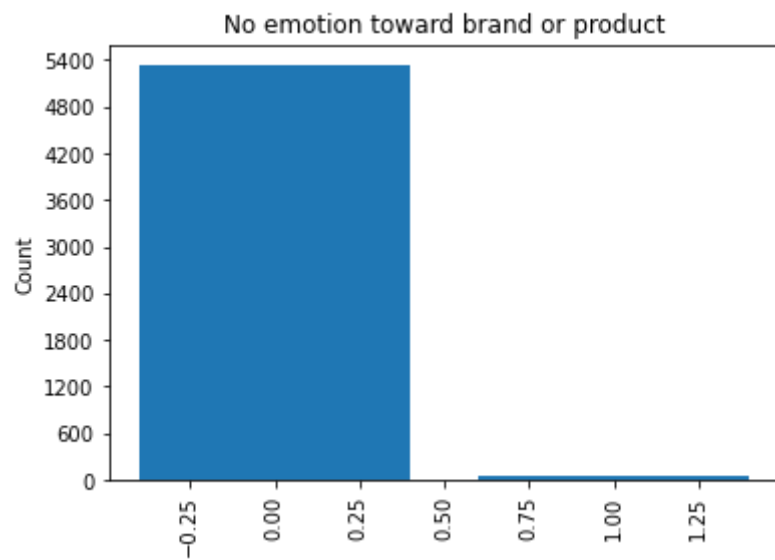
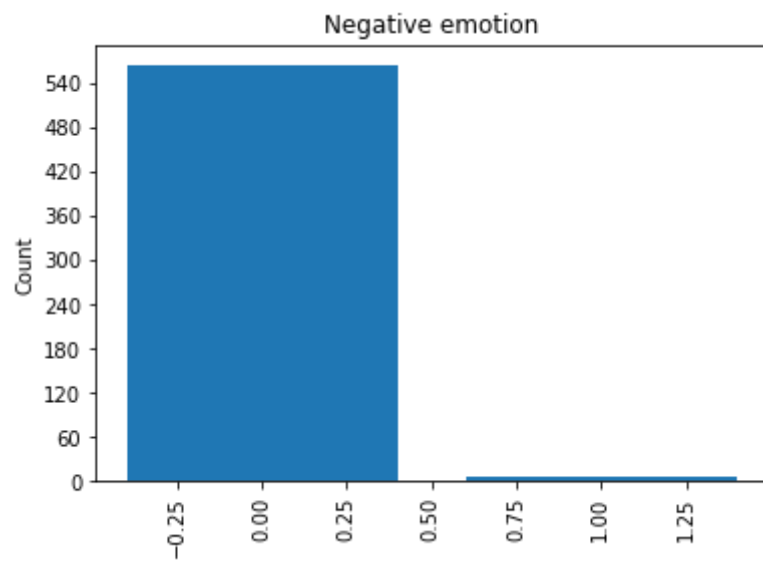
df["contains_emoticon"] = df["tweet"].str.contains(emoticon_query)

def freq_distribution_by_emotion(df):
    emotions = set(df['emotion'].values.tolist())

    #in the loop, create a freq dist for only a subject, then call the visualize

    for em in emotions:
        emotion_df = df[df['emotion'] == em]
        emotion_df_freq_dist = FreqDist(emotion_df["contains_emoticon"].explode())
        visualize_top_10(emotion_df_freq_dist, em)

freq_distribution_by_emotion(df)
```



```
In [44]: X_train["contains_emoticon"] = X_train["tweet"].str.contains(emoticon_query)
X_train.head()
```

| | tweet | num_sentences | contains_emoticon |
|------|---|---------------|-------------------|
| 3363 | having a great time at the google party #sxswi... | 1 | False |
| 3204 | @mention from:ubersocial for iphone now in the... | 1 | False |
| 4460 | are you in town for #sxsw? be sure to check in... | 3 | False |
| 2311 | the ironic tee has been usurped by the ipad 2 ... | 1 | False |
| 6298 | rt @mention marissa mayer: google will connect... | 1 | False |

```
In [45]: X_train.contains_emoticon.value_counts()
```

```
Out[45]: False    6600
         True     102
         Name: contains_emoticon, dtype: int64
```

```
In [46]: X_train.shape
```

```
Out[46]: (6702, 3)
```

```
In [47]: tfidf = TfidfVectorizer(
          max_features=10,
          stop_words=stopwords_list,
        )

# Fit the vectorizer on X_train["text"] and transform it
X_train_vectorized = tfidf.fit_transform(X_train['tweet'])

# Create a full df of vectorized + engineered features
X_train_vectorized_df = pd.DataFrame(X_train_vectorized.toarray(), columns=tfidf.get_feature_names_out())

preprocessed_X_train = pd.concat([X_train_vectorized_df, X_train.reset_index()[['num_sentences', 'contains_emoticon']],
                                  axis=1)

preprocessed_X_train
```

| | app | apple | austin | circles | google | ipad | iphone | new | social | store | num_ |
|------|----------|----------|----------|---------|--------|------|----------|-----|--------|----------|------|
| 0 | 0.000000 | 0.000000 | 0.000000 | 0.0 | 1.0 | 0.0 | 0.000000 | 0.0 | 0.0 | 0.000000 | |
| 1 | 0.656654 | 0.000000 | 0.000000 | 0.0 | 0.0 | 0.0 | 0.52949 | 0.0 | 0.0 | 0.537072 | |
| 2 | 0.000000 | 0.000000 | 0.000000 | 0.0 | 1.0 | 0.0 | 0.000000 | 0.0 | 0.0 | 0.000000 | |
| 3 | 0.000000 | 0.000000 | 0.000000 | 0.0 | 0.0 | 1.0 | 0.000000 | 0.0 | 0.0 | 0.000000 | |
| 4 | 0.000000 | 0.000000 | 0.000000 | 0.0 | 1.0 | 0.0 | 0.000000 | 0.0 | 0.0 | 0.000000 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 6697 | 0.000000 | 0.000000 | 0.000000 | 0.0 | 1.0 | 0.0 | 0.000000 | 0.0 | 0.0 | 0.000000 | |
| 6698 | 0.000000 | 0.000000 | 0.000000 | 0.0 | 0.0 | 0.0 | 1.000000 | 0.0 | 0.0 | 0.000000 | |
| 6699 | 0.000000 | 0.000000 | 0.000000 | 0.0 | 0.0 | 0.0 | 0.000000 | 0.0 | 0.0 | 0.000000 | |
| 6700 | 0.000000 | 0.000000 | 0.000000 | 0.0 | 0.0 | 1.0 | 0.000000 | 0.0 | 0.0 | 0.000000 | |
| 6701 | 0.000000 | 0.75127 | 0.500424 | 0.0 | 0.0 | 0.0 | 0.000000 | 0.0 | 0.0 | 0.430313 | |

6702 rows x 12 columns

```
In [48]: #Score the new model
preprocessed_cv = cross_val_score(baseline_model, preprocessed_X_train, y_train[preprocessed_cv])
```

Out[48]: array([0.60850112, 0.61148397, 0.60895522, 0.60447761, 0.59776119])

```
In [49]: print(baseline_cv.mean())
print(stopwords_removed_cv.mean())
print(stemmed_cv.mean())
print(preprocessed_cv.mean())
```

0.6016115173572851
0.6056398098990506
0.6029537992364797
0.6062358231215288

The Preprocessed is the best version but it's still incredibly marginal. How can I increase it?

```
In [50]: #Trying more max features

tfidf = TfidfVectorizer(
    max_features=200,
    stop_words=stopwords_list,
)

# Fit the vectorizer on X_train["text"] and transform it
X_train_vectorized = tfidf.fit_transform(X_train['tweet'])

# Create a full df of vectorized + engineered features
X_train_vectorized_df = pd.DataFrame(X_train_vectorized.toarray(), columns=tfidf.get_feature_names())

final_X_train = pd.concat([X_train_vectorized_df, X_train.reset_index()[["num_seconds"]]], axis=1)

final_X_train
```

Out[50]:

| | 10 | 11 | 2011 | 30 | 6th | already | also | android | anyone | app | ... | win | winning |
|------|-----|-----|------|-----|-----|---------|------|---------|--------|----------|-----|----------|---------|
| 0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.000000 | ... | 0.000000 | 0.0 |
| 1 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.300485 | ... | 0.000000 | 0.0 |
| 2 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.000000 | ... | 0.000000 | 0.0 |
| 3 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.000000 | ... | 0.000000 | 0.0 |
| 4 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.000000 | ... | 0.000000 | 0.0 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 6697 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.000000 | ... | 0.000000 | 0.0 |
| 6698 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.000000 | ... | 0.000000 | 0.0 |
| 6699 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.000000 | ... | 0.000000 | 0.0 |
| 6700 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.000000 | ... | 0.461711 | 0.0 |
| 6701 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.000000 | ... | 0.000000 | 0.0 |

6702 rows x 202 columns


```
Out[57]:
```


| | 10 | 11 | 2011 | 30 | 6th | already | also | android | anyone | app | ... | win | winning |
|------|-----|-----|----------|-----|-----|---------|------|----------|----------|----------|-----|-----|---------|
| 0 | 0.0 | 0.0 | 0.000000 | 0.0 | 0.0 | 0.0 | 0.0 | 0.000000 | 0.000000 | 0.000000 | ... | 0.0 | 0.0 |
| 1 | 0.0 | 0.0 | 0.000000 | 0.0 | 0.0 | 0.0 | 0.0 | 0.000000 | 0.000000 | 0.000000 | ... | 0.0 | 0.0 |
| 2 | 0.0 | 0.0 | 0.000000 | 0.0 | 0.0 | 0.0 | 0.0 | 0.000000 | 0.000000 | 0.000000 | ... | 0.0 | 0.0 |
| 3 | 0.0 | 0.0 | 0.589882 | 0.0 | 0.0 | 0.0 | 0.0 | 0.000000 | 0.000000 | 0.000000 | ... | 0.0 | 0.0 |
| 4 | 0.0 | 0.0 | 0.000000 | 0.0 | 0.0 | 0.0 | 0.0 | 0.306804 | 0.000000 | 0.000000 | ... | 0.0 | 0.0 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 2229 | 0.0 | 0.0 | 0.000000 | 0.0 | 0.0 | 0.0 | 0.0 | 0.000000 | 0.000000 | 0.000000 | ... | 0.0 | 0.0 |
| 2230 | 0.0 | 0.0 | 0.000000 | 0.0 | 0.0 | 0.0 | 0.0 | 0.000000 | 0.000000 | 0.000000 | ... | 0.0 | 0.0 |
| 2231 | 0.0 | 0.0 | 0.000000 | 0.0 | 0.0 | 0.0 | 0.0 | 0.000000 | 0.778426 | 0.000000 | ... | 0.0 | 0.0 |
| 2232 | 0.0 | 0.0 | 0.000000 | 0.0 | 0.0 | 0.0 | 0.0 | 0.000000 | 0.000000 | 0.831881 | ... | 0.0 | 0.0 |
| 2233 | 0.0 | 0.0 | 0.000000 | 0.0 | 0.0 | 0.0 | 0.0 | 0.000000 | 0.000000 | 0.000000 | ... | 0.0 | 0.0 |

2234 rows x 202 columns

```
In [58]: #final model test score
final_model.score(final_X_test, y_test)
```

Out[58]: 0.6459265890778872

```
In [59]: #Plot confusion matrix
from sklearn.metrics import plot_confusion_matrix
plot_confusion_matrix(final_model, final_X_test, y_test);
```




```
In [60]: # F1 Score of svc model
svc_preds = final_model.predict(final_X_test)
scv_F1 = f1_score(y_test, svc_preds, average='weighted')
scv_F1
```

Out[60]: 0.5845216199204728

This model can't predict negative labels for some reason, or it just learned not to? But its still doing better than random guessing of 'No Emotion' which is kind of impressive.

MultinomialNB Confusion Matrix

```
In [61]: baseline_model.fit(final_X_train, y_train)
         plot_confusion_matrix(baseline_model, final_X_test, y_test);
```



```
In [62]: #MultinomialNB F1 score
         multiNB_preds = baseline_model.predict(final_X_test)
         multiBN_f1 = f1_score(y_test, multiNB_preds, average='weighted')
         multiBN_f1
```

Out[62]: 0.5796038842412904

Logistic Regression and Random Forests

```
In [63]: # Use pipelines to create the models
         rf = Pipeline([('Random Forest', RandomForestClassifier(n_estimators=100, verbose
         lr = Pipeline([('Logistic Regression', LogisticRegression()))
         xgb = Pipeline([('XGB Clf', XGBClassifier())])
```

```
In [64]: # Create list of models
         models = [('Random Forest', rf),
                   ('Logistic Regression', lr),
                   ('XGBoost', xgb)]
```

```
In [65]: # Train the models and obtain the scores
         scores = [(name, cross_val_score(model, final_X_train, y_train, cv=5).mean()) for
         scores
```

```
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 100 out of 100 | elapsed: 1.5s finished
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 100 out of 100 | elapsed: 0.0s finished
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 100 out of 100 | elapsed: 1.5s finished
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
```

```
[Parallel(n_jobs=1)]: Done 100 out of 100 | elapsed:    0.0s finished
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 100 out of 100 | elapsed:    1.5s finished
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 100 out of 100 | elapsed:    0.0s finished
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 100 out of 100 | elapsed:    1.5s finished
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 100 out of 100 | elapsed:    0.0s finished
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 100 out of 100 | elapsed:    1.5s finished
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 100 out of 100 | elapsed:    0.0s finished
```

```
Out[65]: [('Random Forest', 0.6590586218794172),
          ('Logistic Regression', 0.650404465368905),
          ('XGBoost', 0.6602528743308068)]
```

Testing sentiment analysis with VADER

```
In [66]: from nltk.sentiment.vader import SentimentIntensityAnalyzer

sid = SentimentIntensityAnalyzer()

#Load the data and preprocess similarly to before

vader_df = pd.read_csv('Data/data.csv', engine='python')
vader_df.columns = ['tweet', 'subject', 'emotion']

vader_df['emotion'] = vader_df['emotion'].replace("I can't tell", "No emotion to
vader_df['emotion'] = vader_df['emotion'].replace("No emotion toward brand or pr

vader_df = vader_df.dropna(subset=['tweet'])

vader_df.head()
```

```
Out[66]:
```

| | tweet | subject | emotion |
|---|---|--------------------|------------------|
| 0 | .@wesley83 I have a 3G iPhone. After 3 hrs twe... | iPhone | Negative emotion |
| 1 | @jessedee Know about @fludapp ? Awesome iPad/i... | iPad or iPhone App | Positive emotion |
| 2 | @swonderlin Can not wait for #iPad 2 also. The... | iPad | Positive emotion |
| 3 | @sxsw I hope this year's festival isn't as cra... | iPad or iPhone App | Negative emotion |
| 4 | @sxtxstate great stuff on Fri #SXSW: Marissa M... | Google | Positive emotion |

```
In [67]: vader_df.emotion.value_counts()
```

```
Out[67]: No emotion      5544
Positive emotion    2978
Negative emotion     570
Name: emotion, dtype: int64
```

```
In [68]: vader_df['scores'] = vader_df['tweet'].apply(lambda tweet: sid.polarity_scores(t
vader_df.head())
```


| Out[68]: | | tweet | subject | emotion | scores |
|----------|--|--------------------|------------------|---|--------|
| 0 | .@wesley83 I have a 3G iPhone. After 3 hrs twe... | iPhone | Negative emotion | {'neg': 0.203, 'neu': 0.797, 'pos': 0.0, 'comp... | |
| 1 | @jessedee Know about @fludapp ? Awesome iPad/i... | iPad or iPhone App | Positive emotion | {'neg': 0.0, 'neu': 0.576, 'pos': 0.424, 'comp... | |
| 2 | @swonderlin Can not wait for #iPad 2 also. The... | iPad | Positive emotion | {'neg': 0.0, 'neu': 1.0, 'pos': 0.0, 'compound... | |
| 3 | @sxsw I hope this year's festival isn't as cra... | iPad or iPhone App | Negative emotion | {'neg': 0.0, 'neu': 0.663, 'pos': 0.337, 'comp... | |
| 4 | @sxtxstate great stuff on Fri #SXSW: Marissa M... | Google | Positive emotion | {'neg': 0.0, 'neu': 0.796, 'pos': 0.204, 'comp... | |

```
In [69]: vader_df['compound'] = vader_df['scores'].apply(lambda score_dict: score_dict['c
vader_df.head()
```

| Out[69]: | | tweet | subject | emotion | scores | compound |
|----------|--|---|--------------------|------------------|---|----------|
| 0 | | .@wesley83 I have a 3G iPhone. After 3 hrs twe... | iPhone | Negative emotion | {'neg': 0.203, 'neu': 0.797, 'pos': 0.0, 'comp... | -0.6800 |
| 1 | | @jessedee Know about @fludapp ? Awesome iPad/i... | iPad or iPhone App | Positive emotion | {'neg': 0.0, 'neu': 0.576, 'pos': 0.424, 'comp... | 0.9100 |
| 2 | | @swonderlin Can not wait for #iPad 2 also. The... | iPad | Positive emotion | {'neg': 0.0, 'neu': 1.0, 'pos': 0.0, 'compound... | 0.0000 |
| 3 | | @sxsw I hope this year's festival isn't as cra... | iPad or iPhone App | Negative emotion | {'neg': 0.0, 'neu': 0.663, 'pos': 0.337, 'comp... | 0.7269 |
| 4 | | @sxtxstate great stuff on Fri #SXSW: Marissa M... | Google | Positive emotion | {'neg': 0.0, 'neu': 0.796, 'pos': 0.204, 'comp... | 0.6249 |

```
In [70]: def emotion_helper(score):
if score >= .333:
return "Positive emotion"
elif score <= -.333:
return "Negative emotion"
else:
return "No emotion"

vader_df['comp_score'] = vader_df['compound'].apply(emotion_helper)

vader_df.head()
```

| Out[70]: | | tweet | subject | emotion | scores | compound | comp_score |
|----------|---|--------------------|------------------|---|---------|------------------|------------|
| 0 | .@wesley83 I have a 3G iPhone. After 3 hrs twe... | iPhone | Negative emotion | {'neg': 0.203, 'neu': 0.797, 'pos': 0.0, 'comp... | -0.6800 | Negative emotion | |
| 1 | @jessedee Know about @fludapp ? Awesome iPad/i... | iPad or iPhone App | Positive emotion | {'neg': 0.0, 'neu': 0.576, 'pos': 0.424, 'comp... | 0.9100 | Positive emotion | |
| 2 | @swonderlin Can not wait for #iPad 2 also. | iPad | Positive emotion | {'neg': 0.0, 'neu': 1.0, 'pos': 0.0, | 0.0000 | No emotion | |

| | tweet | subject | emotion | scores | compound | comp_score |
|---|---|--------------------|------------------|---|----------|------------------|
| | The... | | | 'compound... | | |
| 3 | @sxsw I hope this year's festival isn't as cra... | iPad or iPhone App | Negative emotion | {'neg': 0.0, 'neu': 0.663, 'pos': 0.337, 'comp... | 0.7269 | Positive emotion |
| 4 | @sxtxstate great stuff on Fri #SXSW: Marissa M... | Google | Positive emotion | {'neg': 0.0, 'neu': 0.796, 'pos': 0.204, 'comp... | 0.6249 | Positive emotion |

```
In [71]: vader_df['Correct'] = vader_df.apply(lambda x: True if x['emotion'] == x['comp_s
```

```
In [72]: vader_df.head()
```

| | tweet | subject | emotion | scores | compound | comp_score | Correct |
|---|---|--------------------|------------------|---|----------|------------------|---------|
| 0 | .@wesley83 I have a 3G iPhone. After 3 hrs twe... | iPhone | Negative emotion | {'neg': 0.203, 'neu': 0.797, 'pos': 0.0, 'comp... | -0.6800 | Negative emotion | True |
| 1 | @jessedee Know about @fludapp ? Awesome iPad/i... | iPad or iPhone App | Positive emotion | {'neg': 0.0, 'neu': 0.576, 'pos': 0.424, 'comp... | 0.9100 | Positive emotion | True |
| 2 | @swonderlin Can not wait for #iPad 2 also. The... | iPad | Positive emotion | {'neg': 0.0, 'neu': 1.0, 'pos': 0.0, 'compound... | 0.0000 | No emotion | False |
| 3 | @sxsw I hope this year's festival isn't as cra... | iPad or iPhone App | Negative emotion | {'neg': 0.0, 'neu': 0.663, 'pos': 0.337, 'comp... | 0.7269 | Positive emotion | False |
| 4 | @sxtxstate great stuff on Fri #SXSW: Marissa M... | Google | Positive emotion | {'neg': 0.0, 'neu': 0.796, 'pos': 0.204, 'comp... | 0.6249 | Positive emotion | True |

```
In [73]: vader_df.Correct.value_counts()
```

```
Out[73]: True      5192
False      3900
Name: Correct, dtype: int64
```

```
In [74]: 5192/(3900+5192)
```

```
Out[74]: 0.5710514738231413
```

This is worse than my initial attempts going to scrap VADER and move onto something else

Trying to use BERT

```
In [75]: #Unused

import torch
from transformers import AutoTokenizer, DistilBertForSequenceClassification
```

```

tokenizer = AutoTokenizer.from_pretrained("distilbert-base-uncased")
model = DistilBertForSequenceClassification.from_pretrained("distilbert-base-uncased")

inputs = tokenizer("I Love my girlfriend!", return_tensors="pt")

with torch.no_grad():
    logits = model(**inputs).logits

predicted_class_ids = torch.arange(0, logits.shape[-1])[torch.sigmoid(logits).squeeze().gt(0.5)]

# To train a model on `num_labels` classes, you can pass `num_labels=num_labels`
num_labels = 3
model = DistilBertForSequenceClassification.from_pretrained(
    "distilbert-base-uncased", num_labels=num_labels, problem_type="multi_label_classification"
)

labels = torch.sum(
    torch.nn.functional.one_hot(predicted_class_ids[None, :].clone(), num_classes=num_labels)
).to(torch.float)
loss = model(**inputs, labels=labels).loss

```

Some weights of the model checkpoint at distilbert-base-uncased were not used when initializing DistilBertForSequenceClassification: ['vocab_layer_norm.weight', 'vocab_transform.bias', 'vocab_projector.weight', 'vocab_transform.weight', 'vocab_projector.bias', 'vocab_layer_norm.bias']

- This IS expected if you are initializing DistilBertForSequenceClassification from the checkpoint of a model trained on another task or with another architecture (e.g. initializing a BertForSequenceClassification model from a BertForPreTraining model).

- This IS NOT expected if you are initializing DistilBertForSequenceClassification from the checkpoint of a model that you expect to be exactly identical (initializing a BertForSequenceClassification model from a BertForSequenceClassification model).

Some weights of DistilBertForSequenceClassification were not initialized from the model checkpoint at distilbert-base-uncased and are newly initialized: ['classifier.bias', 'classifier.weight', 'pre_classifier.weight', 'pre_classifier.bias']

You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

Some weights of the model checkpoint at distilbert-base-uncased were not used when initializing DistilBertForSequenceClassification: ['vocab_layer_norm.weight', 'vocab_transform.bias', 'vocab_projector.weight', 'vocab_transform.weight', 'vocab_projector.bias', 'vocab_layer_norm.bias']

- This IS expected if you are initializing DistilBertForSequenceClassification from the checkpoint of a model trained on another task or with another architecture (e.g. initializing a BertForSequenceClassification model from a BertForPreTraining model).

- This IS NOT expected if you are initializing DistilBertForSequenceClassification from the checkpoint of a model that you expect to be exactly identical (initializing a BertForSequenceClassification model from a BertForSequenceClassification model).

Some weights of DistilBertForSequenceClassification were not initialized from the model checkpoint at distilbert-base-uncased and are newly initialized: ['classifier.bias', 'classifier.weight', 'pre_classifier.weight', 'pre_classifier.bias']

You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

Neural Net Attempt

```
In [76]: # Neural Net imports

from keras.preprocessing.sequence import pad_sequences
from keras.layers import Input, Dense, LSTM, Embedding
from keras.layers import Dropout, Activation, Bidirectional, GlobalMaxPool1D
from keras.models import Sequential
from keras import initializers, regularizers, constraints, optimizers, layers
from keras.preprocessing import text, sequence
```

```
In [77]: # Same data cleaning from the start, but shorter

nn_df = pd.read_csv('Data/data.csv', engine='python')
nn_df.columns = ['tweet', 'subject', 'emotion']
nn_df = nn_df.drop(unknowns_df.index.tolist())
nn_df = nn_df.dropna(subset=['tweet'])
nn_df['tweet'] = nn_df['tweet'].str.lower()
```

Now that the data is cleaned again, I need to preprocess is for Keras. That means creating sequences of the same length for the Neural Net to read.

```
In [78]: # Create a tokenizer using the 20,000 most used words
# Fit on the tweets
# Create a list of the tweets
# Set the list of tweets as your training data and pad them for even input into

tokenizer = text.Tokenizer(num_words=20000)
tokenizer.fit_on_texts(list(nn_df['tweet']))
list_tokenized_tweets = tokenizer.texts_to_sequences(df['tweet'])
X_t = sequence.pad_sequences(list_tokenized_tweets, maxlen=140)
```

Set Y values

```
In [79]: #Set target (y) value as emotion
target = nn_df['emotion']
y = pd.get_dummies(target).values
```

Create the model.

The model will consist of an initial layer with an embedding size of 20,000 because we used a tokenizer with the 20,000 most used words. A dropout layer with relu activation. And I final layer with 3 classes and softmax, these are because its a multi-category classification problem.

```
In [80]: # Create the model and layers
# Text models should use an embedding layer as the start
# Final layer is 3 for the output size (3 classes)

model = Sequential()

embedding_size = 128
model.add(Embedding(20000, embedding_size))
model.add(LSTM(25, return_sequences=True))
model.add(GlobalMaxPool1D())
model.add(Dropout(0.5))
```

```
model.add(Dense(50, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(3, activation='softmax'))
```

Don't forget to compile

```
In [81]: # Compile Model
# Use 'categorical_crossentropy' for multi category classification
model.compile(loss='categorical_crossentropy',
              optimizer='adam',
              metrics=['accuracy'])
```

```
In [82]: # Summary print out
model.summary()
```

Model: "sequential"

| Layer (type) | Output Shape | Param # |
|------------------------------|-------------------|---------|
| embedding (Embedding) | (None, None, 128) | 2560000 |
| lstm (LSTM) | (None, None, 25) | 15400 |
| global_max_pooling1d (Global | (None, 25) | 0 |
| dropout (Dropout) | (None, 25) | 0 |
| dense (Dense) | (None, 50) | 1300 |
| dropout_1 (Dropout) | (None, 50) | 0 |
| dense_1 (Dense) | (None, 3) | 153 |
| Total params: 2,576,853 | | |
| Trainable params: 2,576,853 | | |
| Non-trainable params: 0 | | |

I chose 3 epochs because I did some testing that showed only 3 were necessary, more on that below

```
In [83]: #Train the model
model.fit(X_t, y, epochs=3, batch_size=32, validation_split=0.1)
```

```
Epoch 1/3
252/252 [=====] - 19s 74ms/step - loss: 0.8926 - accuracy: 0.5856 - val_loss: 0.7829 - val_accuracy: 0.6107
Epoch 2/3
252/252 [=====] - 16s 64ms/step - loss: 0.7592 - accuracy: 0.6677 - val_loss: 0.7007 - val_accuracy: 0.6577
Epoch 3/3
252/252 [=====] - 16s 65ms/step - loss: 0.6083 - accuracy: 0.7538 - val_loss: 0.7023 - val_accuracy: 0.6711
```

```
Out[83]: <tensorflow.python.keras.callbacks.History at 0x7f915dcae040>
```

It looks like 3 epochs is actually the sweet spot and anything past that gives diminishing results. At 3 epochs we see .85 accuracy on the training data and .66 accuracy on the validation set. The training accuracy is much higher than our simpler models and the validation accuracy is right around where the training accuracy was so I would say this is a successful test.

```
In [84]: keras_preds = model.predict_classes(X_t)
```

WARNING:tensorflow:From <ipython-input-84-e192ed4d5f15>:1: Sequential.predict_classes (from tensorflow.python.keras.engine.sequential) is deprecated and will be removed after 2021-01-01.

Instructions for updating:

Please use instead: * `np.argmax(model.predict(x), axis=-1)`, if your model does multi-class classification (e.g. if it uses a `softmax` last-layer activation) * `(model.predict(x) > 0.5).astype("int32")`, if your model does binary classification (e.g. if it uses a `sigmoid` last-layer activation).

```
In [85]: keras_preds
```

```
Out[85]: array([0, 2, 2, ..., 1, 1, 1])
```

```
In [86]: target = pd.DataFrame(y)
target.columns = [0, 1, 2]
target = target.stack()
target = pd.Categorical(target[target!=0].index.get_level_values(1))
target = pd.DataFrame(target)
target.columns = ['Emotion']
```

```
In [87]: # Print Confusion Matrix
print(confusion_matrix(target, keras_preds))
```


```
[[ 207  234  129]
 [   30 5094  264]
 [   14  996 1968]]
```

```
In [133... # Import ConfusionMatrixDisplay to display the CM right above this cell, but pre

from sklearn.metrics import ConfusionMatrixDisplay

disp = ConfusionMatrixDisplay(confusion_matrix(target, keras_preds))
disp.plot()
```

```
Out[133... <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x7f915d1c94f0
>
```



```
In [88]: # Create Validation set
# Make X_t a dataframe
X_full = pd.DataFrame(X_t)

# Create a random sample of 700 rows from X
X_val = X_full.sample(n=1800, random_state=42)
```

```
# Create a subset of Y with the same indices as X_subset
Y_val = target.loc[X_val.index]
```

```
In [89]: # Predict on validation set
val_preds = model.predict_classes(X_val)
```

```
In [90]: # Keras F1 Score
keras_f1 = f1_score(Y_val, val_preds, average='weighted')
keras_f1
```

```
Out[90]: 0.8023441831895548
```

Not too much of a surprise that the neural net has the best F1 score so far. This is our best model at the moment in both accuracy and F1

Some Additional EDA

```
In [91]: df.head()
```


```
Out[91]:
```

| | tweet | subject | emotion | tweet_tokenized | tweet_without_stopwords | num_sentences | col |
|---|---|--------------------------|---------------------|---|--|---------------|-----|
| 0 | .@wesley83 i have a 3g iphone. after 3 hrs twe... | iPhone | Negative emotion | [wesley83, have, 3g, iphone, after, hrs, tweet... | [wesley83, 3g, iphone, hrs, tweeting, rise_aus... | 5 | |
| 1 | @jessedee know about @fludapp ? awesome ipad/i... | iPad or iPhone App | Positive emotion | [jessedee, know, about, fludapp, awesome, ipad... | [jessedee, know, fludapp, awesome, ipad, iphon... | 3 | |
| 2 | @swonderlin can not wait for #ipad 2 also. the... | iPad | Positive emotion | [swonderlin, can, not, wait, for, ipad, also, ... | [swonderlin, wait, ipad, also, sale] | 2 | |
| 3 | @sxsw i hope this year's festival isn't as cra... | iPad or iPhone App | Negative emotion | [sxsw, hope, this, year, festival, isn, as, cr... | [hope, year, festival, crashy, year, iphone, app] | 2 | |
| 4 | @sxtxstate great stuff on fri #sxsw: marissa m... | Google | Positive emotion | [sxtxstate, great, stuff, on, fri, sxsw, maris... | [sxtxstate, great, stuff, fri, marissa, mayer,... | 1 | |

```
In [92]: plt.hist(nn_df['tweet'].str.len(), bins=20, label='train')
plt.legend()
plt.show()
```



```
plt.axis('off')
plt.show()
```



In [128...

```
apple_words = ' '.join([text for text in df['tidy_tweets'][df['subject'] == 'Apple'])

wordcloud = WordCloud(width=800, height=500, random_state=21, max_font_size=110)
plt.figure(figsize=(10, 7))
plt.imshow(wordcloud, interpolation="bilinear")
plt.axis('off')
plt.show()
```



In [98]:

```
tokenized_tweet = df['tidy_tweets'].apply(lambda x: x.split()) # tokenizing
```

```

import gensim

model_w2v = gensim.models.Word2Vec(
    tokenized_tweet,
    size=200, # desired no. of features/independent variables
    window=5, # context window size
    min_count=2, # Ignores all words with total frequency lower than 2.
    sg = 1, # 1 for skip-gram model
    hs = 0,
    negative = 10, # for negative sampling
    workers= 32, # no.of cores
    seed = 34
)

model_w2v.train(tokenized_tweet, total_examples=len(df['tidy_tweets']), epochs=

```

Out[98]: (1320662, 1671440)

In [99]: df['tidy_tweets']

```

Out[99]: 0    wesley83 3g iphone hrs tweeting rise_austin de...
1    jessedee know fludapp awesome ipad iphone app ...
2                                swonderlin wait ipad also sale
3                hope year festival crashy year iphone app
4    sctxstate great stuff fri marissa mayer google...
...
9088                                ipad everywhere
9089    wave buzz interrupt regularly scheduled geek p...
9090    google zeiger physician never reported potenti...
9091    verizon iphone customers complained time fell ...
9092                                google tests check offers
Name: tidy_tweets, Length: 8936, dtype: object

```

```

In [100... total_vocabulary = []

for tweet in df['tweet_without_stopwords']:
    for word in tweet:
        total_vocabulary.append(word)

total_vocabulary = set(total_vocabulary)

```

```

In [101... len(total_vocabulary)
print('There are {} unique tokens in the dataset.'.format(len(total_vocabulary)))

There are 9453 unique tokens in the dataset.

```

```

In [102... glove = {}
with open('glove.twitter.27B.50d.txt', 'rb') as f:
    for line in f:
        parts = line.split()
        word = parts[0].decode('utf-8')
        if word in total_vocabulary:
            vector = np.array(parts[1:], dtype=np.float32)
            glove[word] = vector

```

```

In [103... class W2vVectorizer(object):

    def __init__(self, w2v):
        # Takes in a dictionary of words and vectors as input
        self.w2v = w2v
        if len(w2v) == 0:

```

```

        self.dimensions = 0
    else:
        self.dimensions = len(w2v[next(iter(glove))])

    # Note: Even though it doesn't do anything, it's required that this object i
    # it can't be used in a scikit-learn pipeline
    def fit(self, X, y):
        return self

    def transform(self, X):
        return np.array([
            np.mean([self.w2v[w] for w in words if w in self.w2v]
                    or [np.zeros(self.dimensions)], axis=0) for words in X])

```

Word2Vec of the previous models

```

In [104... xgb = Pipeline([('Word2Vec Vectorizer', W2vVectorizer(glove)),
                  ('XGB Clf', XGBClassifier())])
rf = Pipeline([('Word2Vec Vectorizer', W2vVectorizer(glove)),
               ('Random Forest', RandomForestClassifier(n_estimators=100, verbose
svc = Pipeline([('Word2Vec Vectorizer', W2vVectorizer(glove)),
               ('Support Vector Machine', SVC())])
lr = Pipeline([('Word2Vec Vectorizer', W2vVectorizer(glove)),
               ('Logistic Regression', LogisticRegression())])

```

```

In [105... models = [('Random Forest', rf),
                  ('Support Vector Machine', svc),
                  ('Logistic Regression', lr),
                  ('XGBoost', xgb)]

```

```

In [106... scores = [(name, cross_val_score(model, df['tweet_without_stopwords'], df['emoti

[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 100 out of 100 | elapsed: 3.6s finished
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 100 out of 100 | elapsed: 0.0s finished
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 100 out of 100 | elapsed: 3.6s finished
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 100 out of 100 | elapsed: 0.0s finished
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 100 out of 100 | elapsed: 3.4s finished
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 100 out of 100 | elapsed: 0.0s finished
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 100 out of 100 | elapsed: 3.2s finished
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 100 out of 100 | elapsed: 0.0s finished
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 100 out of 100 | elapsed: 3.2s finished
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 100 out of 100 | elapsed: 0.0s finished

```

```

In [107... scores

```

```

Out[107... [('Random Forest', 0.6820725498222935),
            ('Support Vector Machine', 0.6497306547786713),
            ('Logistic Regression', 0.6337273047075009),
            ('XGBoost', 0.6827426266510932)]

```

In [143...

```
#Fit, Predict, F1

rf.fit(df['tweet_without_stopwords'], df['emotion'])

rf_preds = rf.predict(df['tweet_without_stopwords'])

rf_f1 = f1_score(df['emotion'], rf_preds, average="weighted")

print("F1 Score: ", rf_f1)
```


```
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 100 out of 100 | elapsed: 4.5s finished
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
F1 Score: 0.9483759872470012
[Parallel(n_jobs=1)]: Done 100 out of 100 | elapsed: 0.1s finished
```

In [144...

```
cm = confusion_matrix(df['emotion'], rf_preds)
disp = ConfusionMatrixDisplay(cm)
disp.plot()
```

Out[144...

```
<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x7f9168bb0b50>
```



In [109...

```
#Fit, Predict, F1

xgb.fit(df['tweet_without_stopwords'], df['emotion'])

xgb_preds = xgb.predict(df['tweet_without_stopwords'])

xgb_f1 = f1_score(df['emotion'], xgb_preds, average="weighted")

print("F1 Score: ", xgb_f1)
```

```
F1 Score: 0.9473074686501557
```

In [111...

```
# Create Validation set

# Create a random sample of 1800 rows from X
X_val = df.sample(n=1800, random_state=42)

# Create a subset of Y with the same indices as X_subset
Y_val = X_val['emotion']
```



```
In [112... rf_preds = rf.predict(X_val['tweet_without_stopwords'])

rf_f1 = f1_score(Y_val, rf_preds, average="weighted")

print("Random Forest F1 Score: ", rf_f1)
```

Random Forest F1 Score: 0.9464691347279439

[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 100 out of 100 | elapsed: 0.0s finished

```
In [113... xgb_preds = xgb.predict(X_val['tweet_without_stopwords'])

xgb_f1 = f1_score(Y_val, xgb_preds, average="weighted")


print("XGB F1 Score: ", xgb_f1)
```

XGB F1 Score: 0.9398811741374118

The F1 score of our random forest is incredible. This is the best model and by far better than the neural net

```
In [141... cm = confusion_matrix(Y_val, rf_preds)
disp = ConfusionMatrixDisplay(cm)
disp.plot()
```

Out[141... <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x7f91687510a0
>



Lets explore similar words and offer google and apple recommendations

```
In [114... model_w2v.wv.most_similar(positive="google")
```

```
Out[114...] [('brother', 0.5848442316055298),
             ('incorrect', 0.5755859613418579),
             ('22sxsx', 0.5551255941390991),
             ('hee', 0.5488891005516052),
             ('pointing', 0.5466156005859375),
             ('nptech', 0.545316219329834),
             ('ne', 0.5447436571121216),
             ('nadjja', 0.5431619882583618),
             ('spokewoman', 0.5382939577102661),
             ('marisa', 0.5367705821990967)]
```

```
In [115...] model_w2v.wv.most_similar(negative=["google"])
```

```
Out[115...] [('co', 0.0047996994107961655),
             ('kindle', 0.0021751169115304947),
             ('winning', 0.0006275922060012817),
             ('16gb', -0.0008450224995613098),
             ('mophie', -0.0037463074550032616),
             ('bar', -0.00474444217979908),
             ('pure', -0.004777221009135246),
             ('screen', -0.004891552031040192),
             ('black', -0.008278627879917622),
             ('pack', -0.01591699942946434)]
```

```
In [116...] model_w2v.wv.most_similar(positive="apple")
```

```
Out[116...] [('sixth', 0.6366628408432007),
             ('5000', 0.6305968165397644),
             ('tcrn', 0.6255877017974854),
             ('shut', 0.6254369616508484),
             ('impulse', 0.624621570110321),
             ('applestore', 0.6233795285224915),
             ('rage', 0.6223534345626831),
             ('toptweets', 0.6208226680755615),
             ('geekfest', 0.6204372048377991),
             ('brian_lam', 0.6194248795509338)]
```

```
In [117...] model_w2v.wv.most_similar(negative=["apple"])
```

```
Out[117...] [('catch', -0.007958658039569855),
             ('twitter', -0.034292034804821014),
             ('rules', -0.03449346497654915),
             ('panel', -0.04481664299964905),
             ('browser', -0.0509837344288826),
             ('nyt', -0.05322345346212387),
             ('start', -0.05585930123925209),
             ('excel', -0.05867619812488556),
             ('need', -0.06025959178805351),
             ('scrape', -0.06044682115316391)]
```

```
In [118...] model_w2v.wv.most_similar(positive="iphone")
```

```
Out[118...] [('desperate', 0.5116559267044067),
             ('sync', 0.49874168634414673),
             ('handy', 0.4877375066280365),
             ('wew', 0.4863353669643402),
             ('fully', 0.4829857647418976),
             ('schedules', 0.4818391799926758),
             ('charts', 0.4782484769821167),
             ('communications', 0.4768584072589874),
             ('merchant', 0.47590404748916626),
             ('hobo', 0.4734790027141571)]
```

```
In [119... model_w2v.wv.most_similar(negative=["iphone"])
```

```
Out[119... [('jobs', 0.029164083302021027),  
            ('guy', 0.01111283153295517),  
            ('street', 0.007235661149024963),  
            ('part', -0.0006067678332328796),  
            ('owners', -0.008246984332799911),  
            ('front', -0.010976498946547508),  
            ('head', -0.016630683094263077),  
            ('steve', -0.018890613690018654),  
            ('miss', -0.019209370017051697),  
            ('sales', -0.019321508705615997)]
```

```
In [120... model_w2v.wv.most_similar(positive="ipad")
```

```
Out[120... [('adam', 0.49525368213653564),  
            ('attn', 0.4878769516944885),  
            ('relic', 0.4720078408718109),  
            ('resist', 0.4668174982070923),  
            ('channels', 0.4613021910190582),  
            ('kenny', 0.46007850766181946),  
            ('smarty', 0.45260828733444214),  
            ('commercial', 0.45218873023986816),  
            ('baby', 0.4521023631095886),  
            ('smileyparty', 0.45200440287590027)]
```

```
In [121... model_w2v.wv.most_similar(negative=["ipad"])
```

```
Out[121... [('system', 0.02862054854631424),  
            ('engine', -0.006649543531239033),  
            ('town', -0.009210258722305298),  
            ('hotpot', -0.009825445711612701),  
            ('info', -0.018528716638684273),  
            ('recommendation', -0.02047230675816536),  
            ('become', -0.027191689237952232),  
            ('directions', -0.027908606454730034),  
            ('nerds', -0.03062290884554386),  
            ('behind', -0.0346079058945179)]
```

```
In [122... model_w2v.wv.most_similar(positive="app")
```

```
Out[122... [('wew', 0.5776689052581787),  
            ('yayrt', 0.5335705280303955),  
            ('forbes', 0.5316839814186096),  
            ('sync', 0.5285147428512573),  
            ('wedig', 0.5284035205841064),  
            ('workspace', 0.5283859968185425),  
            ('lightbox', 0.5249691009521484),  
            ('concertgoers', 0.5216178894042969),  
            ('nicely', 0.5132554173469543),  
            ('casa', 0.5123087763786316)]
```

```
In [123... model_w2v.wv.most_similar(negative=["app"])
```

```
Out[123... [('30', 0.008424434810876846),  
            ('front', 0.0036735422909259796),  
            ('headline', -0.005456060171127319),  
            ('room', -0.006305336952209473),  
            ('guy', -0.007032092660665512),  
            ('recipe', -0.015117660164833069),  
            ('matt', -0.016400093212723732),  
            ('instead', -0.017644822597503662),
```



```
('existence', -0.02015220746397972),  
( 'pay', -0.020419329404830933) ]
```

Recommendations for Apple

1. One of the most positive words associated with Apple is Store / applestore. People love the Apple Store, it was revolutionary when first introduced. How can you market the store and highlight how great it is, and get people to come in?
2. Some common Positive iphone words are handy and schedules. People love how great the iphone is as a personal device for day to day tasks. Create marketing to highlight businessmen using the iphone for scheduling meetings, calender apointments, etc. Or families scheduling playdates for children and soccer games etc.
3. A recurring negative word is sales. Prices are high and there aren't enough sales for iphones, ipads, etc.

In []: